

P2

February 18, 2021

1 P2 - Team Jarlsberg

1.0.1 Zachary Noel

1.0.2 Siddhesh Mahadeshwar

1.0.3 Erin Dolson

```
[45]: import scipy.spatial.distance as dist
import pandas as pd
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sp
import scipy.stats as stats
import math
import plotly.express as px
from sklearn import metrics
from sklearn.metrics.pairwise import cosine_similarity
from sklearn import cluster
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from scipy.spatial import minkowski_distance
from scipy.spatial.distance import cityblock
from scipy.spatial import distance
from scipy.spatial.distance import cdist
from scipy.spatial import distance_matrix
from gapstat import GapStatClustering
from scipy.cluster import hierarchy
from sklearn.datasets import load_iris
from sklearn import preprocessing
%matplotlib inline
```

2 Part B

2.1 Problem 4

Min-max normalization with min = 0 and max = 1

```
[46]: # Normalize between (0,1)
df = pd.DataFrame([20, 30, 40, 60, 120])
x = df.values.reshape(-1, 1)
y = MinMaxScaler()
x_scaled = y.fit_transform(x)

print(x_scaled)
```

```
[[0. ]
 [0.1]
 [0.2]
 [0.4]
 [1. ]]
```

Min-max normalization with min = -1 and max = 1

```
[47]: # Normalize between (-1,1)
df = pd.DataFrame([20, 30, 40, 60, 120])
x = df.values.reshape(-1, 1)
z = MinMaxScaler(feature_range = (-1, 1))
x_scaled1 = z.fit_transform(x)

print(x_scaled1)
```

```
[[ -1. ]
 [-0.8]
 [-0.6]
 [-0.2]
 [ 1. ]]
```

Z-score normalization

```
[48]: # ZScore Calculation
data = pd.DataFrame([20, 30, 40, 60, 120])
scaler = StandardScaler()

print(scaler.fit_transform(data))
```

```
[[ -0.95632472]
 [-0.67505274]
 [-0.39378077]
 [ 0.16876319]
 [ 1.85639504]]
```

2.2 Problem 5

5A - Calculate and present the distance between the new data point and each of the points in the data set using Manhattan distance, Euclidean distance, Minkowski distance(= 3), supremum distance, and cosine similarity.

```
[49]: # Initialize Data
x1 = [1.4, 1.3, 2.9]
x2 = [1.8, 1.1, 3.2]
x3 = [1.3, 1.2, 2.9]
x4 = [0.9, 3.3, 3.1]
x5 = [1.5, 2.1, 3.3]
x = [1.25, 1.74, 3.01]

df1 = pd.DataFrame((np.array([x1,x2,x3,x4,x5])),
                    columns=['A', 'B', 'C'])
df2 = pd.DataFrame((np.array([x])),
                    columns=['A', 'B', 'C'])

# Manhattan Distance
df1['Manhattan'] = cdist(df1[['A','B','C']].values, df2[['A','B','C']].values,
                        'cityblock')
print(df1['Manhattan'])
print('-----')

# Euclidean Distance
df1['Euclidean'] = cdist(df1[['A','B','C']].values,
                        df2[['A','B','C']].values, 'euclidean')
print(round(df1['Euclidean'], 5))
print('-----')

#Minkowski Distance
df1['Minkowski'] = cdist(df1[['A','B','C']].values, df2[['A','B','C']].values,
                        'minkowski', 3)
print(round(df1['Minkowski'], 5))
print('-----')

# Supremum Distance
df1['Supremum'] = cdist(df1[['A','B','C']].values,
                        df2[['A','B','C']].values, 'chebyshev')
print(df1['Supremum'])
print('-----')

# Cosine Similarity
df1['Cosine Similarity'] = cosine_similarity(df1[['A','B','C']].values,
                        df2[['A','B','C']].values)
print(round(df1['Cosine Similarity'], 5))
print('-----')
```

```

0    0.70
1    1.38
2    0.70
3    2.00
4    0.90
Name: Manhattan, dtype: float64
-----
0    0.47770
1    0.86499
2    0.55335
3    1.60131
4    0.52555
Name: Euclidean, dtype: float64
-----
0    0.44796
1    0.75792
2    0.54166
3    1.56595
4    0.44254
Name: Minkowski, dtype: float64
-----
0    0.44
1    0.64
2    0.54
3    1.56
4    0.36
Name: Supremum, dtype: float64
-----
0    0.99303
1    0.97426
2    0.99133
3    0.94973
4    0.99898
Name: Cosine Similarity, dtype: float64
-----

```

5B - Normalize the data using min-max normalization to be between 0 and 1. What is the Euclidean distance between the new data point and x_1, \dots, x_5 .

```

[50]: # Initialize Data
x1 = [1.4, 1.3, 2.9]
x2 = [1.8, 1.1, 3.2]
x3 = [1.3, 1.2, 2.9]
x4 = [0.9, 3.3, 3.1]
x5 = [1.5, 2.1, 3.3]
x = [1.25, 1.74, 3.01]

df3 = pd.DataFrame((np.array([x1,x2,x3,x4,x5,x])),

```

```

        columns=['A', 'B', 'C'])

print(df3)
print('-----')

# Normalizing the dataframe
df = pd.DataFrame(df3[['A', 'B', 'C']])
y = MinMaxScaler()
df3[['A', 'B', 'C']] = y.fit_transform(df3[['A', 'B', 'C']])
print(round(df3[['A', 'B', 'C']],5))
print('-----')

# Dataframes manipulation
df4 = pd.DataFrame(df3.iloc[5,0:6].values)
df4 = df3
df3 = df3.drop([5])
df4 = df4.drop([0,1,2,3,4])

# Euclidean Distance Calculation with Normalized Dataframes
df3['Euclidean Distance'] = cdist(df3[['A', 'B', 'C']].values,
                                   df4[['A', 'B', 'C']].values, 'euclidean')
print(round(df3['Euclidean Distance'], 5))

```

	A	B	C
0	1.40	1.30	2.90
1	1.80	1.10	3.20
2	1.30	1.20	2.90
3	0.90	3.30	3.10
4	1.50	2.10	3.30
5	1.25	1.74	3.01

	A	B	C
0	0.55556	0.09091	0.000
1	1.00000	0.00000	0.750
2	0.44444	0.04545	0.000
3	0.00000	1.00000	0.500
4	0.66667	0.45455	1.000
5	0.38889	0.29091	0.275

0	0.37869
1	0.82687
2	0.37277
3	0.83945
4	0.79345

Name: Euclidean Distance, dtype: float64

3 Problem 6

```
[51]: info = pd.read_csv("Pokemon.csv",
                        engine='python',)
info2 = pd.read_csv("Pokemon.csv",
                    engine='python',)
info.head()
```

```
[51]:
```

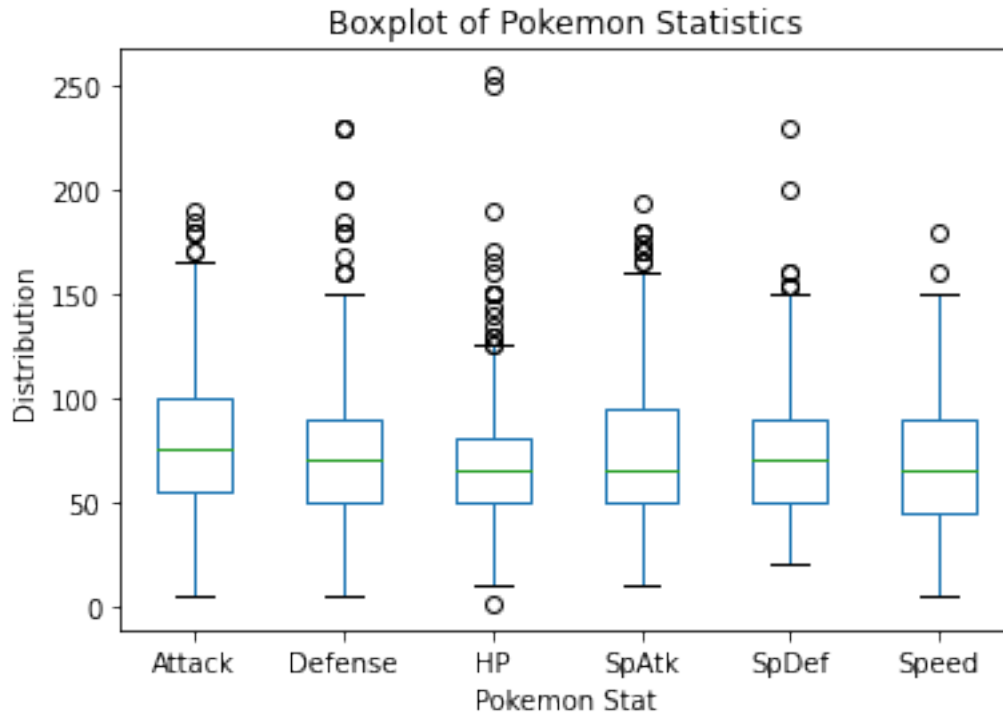
	Num	Name	Type1	...	Speed	Generation	Legendary
0	1	Bulbasaur	Grass	...	45	1	False
1	2	Ivysaur	Grass	...	60	1	False
2	3	Venusaur	Grass	...	80	1	False
3	3	VenusaurMega	Venusaur	Grass	...	1	False
4	4	Charmander	Fire	...	65	1	False

[5 rows x 13 columns]

6A - After loading in the data, look at the distribution of Pokemon features we will use for clustering: HP, Attack, Defense, SpAtk, SpDef, and Speed

```
[52]: fig = plt.figure(figsize=(10,15))
# I looked up how to filter out columns for boxplots and found this at https://
→stackoverflow.com/questions/13003051/
→how-do-i-exclude-a-few-columns-from-a-dataframe-plot
exclude = ['Num', 'Name', 'Total', 'Generation', 'Legendary']
info.loc[:, info.columns.difference(exclude)].plot.box()
plt.ylabel("Distribution")
plt.xlabel("Pokemon Stat")
plt.title("Boxplot of Pokemon Statistics");
```

<Figure size 720x1080 with 0 Axes>



6B - The features have different ranges, therefore we should scale the data before considering the clustering analysis. Scale the data using min-max normalization with range of [0, 1].

```
[53]: exclude = ['Num', 'Name', 'Total', 'Type1', 'Type2', 'Generation', 'Legendary']
y = MinMaxScaler()
info.loc[:, info.columns.difference(exclude)] = y.fit_transform(info.loc[:,
→info.columns.difference(exclude)])
```

6C - Run Kmeans clustering on the data of b with k = [2, 3, ..., 8]

```
[54]: kmtrace = []
for i in range(2, 9):
    km = cluster.KMeans(n_clusters=i)
    km.fit(info.loc[:, info.columns.difference(exclude)])
    kmtrace.append(km)
```

6D - Determine the "best" number of clusters using gap statistic

```
[55]: # utilized source at https://github.com/jmmaloney3/gapstat/blob/master/
→notebooks/GapStatClustering.ipynb to determine this
gstat_ac = GapStatClustering(max_k=8).fit(info.loc[:, info.columns.
→difference(exclude)])
num = gstat_ac.n_clusters_
print(num)
```

6E - Report the mean skill values (centers) of each group, best number of groups determined in (d), as a table/data frame

```
[56]: km = cluster.KMeans(n_clusters=num)
km.fit(info.loc[:, info.columns.difference(exclude)])
km.cluster_centers_
df2 = pd.DataFrame(km.cluster_centers_, columns=['HP', 'Attack', 'Defense',
→'SpAtk', 'SpDef', 'Speed'])
print(df2)
```

	HP	Attack	Defense	SpAtk	SpDef	Speed
0	0.566944	0.427863	0.311387	0.269523	0.253443	0.324747
1	0.624046	0.380228	0.348701	0.621445	0.362483	0.539706
2	0.261892	0.209032	0.201631	0.201262	0.140085	0.251286
3	0.403934	0.260450	0.271928	0.388793	0.246775	0.526184
4	0.339027	0.390293	0.302205	0.430522	0.382210	0.286446

6F - Report the mean skill values (center) of each group, best number of groups determined in (d) as a table/data frame using original data scaling (reverse the scaling back to the original data range)

```
[57]: km3 = cluster.KMeans(n_clusters=num)
km3.fit(info2.loc[:, info.columns.difference(exclude)])
km3.cluster_centers_
df = pd.DataFrame(km3.cluster_centers_, columns=['HP', 'Attack', 'Defense',
→'SpAtk', 'SpDef', 'Speed'])
print(df)
```

	HP	Attack	Defense	SpAtk	SpDef	Speed
0	120.925926	92.879630	89.509259	123.111111	97.861111	100.287037
1	53.753623	51.550725	49.887681	47.322464	48.858696	49.605072
2	96.064706	69.276471	74.670588	69.070588	68.294118	88.835294
3	67.606250	74.512500	83.037500	92.993750	89.700000	67.343750
4	94.848837	129.255814	69.662791	61.372093	87.279070	49.104651

6G - Create a single figure with a radar plot showing the mean skill values (Center) for each cluster group, set from (d)

```
[58]: for i in range(0, num):
    Cluster = pd.DataFrame([df2.iloc[i, 0:6].values],
→columns=(['HP', 'Attack', 'Defense', 'SpAtk', 'SpDef', 'Speed']))
    my_title = "Cluster " + str(i+1)
    fig = px.line_polar(Cluster, r=Cluster.iloc[0, 0:6],
→theta=['HP', 'Attack', 'Defense', 'SpAtk', 'SpDef', 'Speed'], line_close=True,
→title=my_title)
    fig.show()
print("NOTICE: The graph does not appear in the PDF. Please see the end of the
→PDF to see the radar plots that were generated. You can also run the .ipynb
→notebook to see them.")
```


NOTICE: The graph does not appear in the PDF. Please see the end of the PDF to see the radar plots that were generated. You can also run the .ipynb notebook to see them.

3.1 Problem 7

```
[59]: # Import the data
df = pd.read_csv("music2.csv",
                 engine='python')

df.head()
```

```
[59]:
```

	Song	Artist	Type	...	LMax	LFEner	LFreq
0	Dancing Queen	Abba	Rock	...	29921	105.92095	59.57379
1	Knowing Me	Abba	Rock	...	27626	102.83616	58.48031
2	Take a Chance	Abba	Rock	...	26372	102.32488	124.59397
3	Mamma Mia	Abba	Rock	...	28898	101.61648	48.76513
4	Lay All You	Abba	Rock	...	27940	100.30076	74.02039

[5 rows x 8 columns]

Standardize the numeric variables

```
[60]: # SOURCE for Standardization: https://www.kaggle.com/discdiver/guide-to-scaling-and-standardizing
x = df['LVar']
stats.zscore(x)
col_names = ['LVar', 'LAve', 'LMax', 'LFEner', 'LFreq']

shrunk_df = df[['LVar', 'LAve', 'LMax', 'LFEner', 'LFreq']].copy()

s_scaler = preprocessing.StandardScaler()
standardized_df = s_scaler.fit_transform(shrunk_df)

standardized_df = pd.DataFrame(standardized_df, columns=col_names)
```

Perform hierarchial clustering two times, with single and complete linkage. Label the clusters by type, and then by artist

```
[61]: types_list = df['Type'].to_list()
artists_list = df['Artist'].to_list()

# single linking 'Type' of music
Z = hierarchy.linkage(standardized_df, 'single')
plt.figure(figsize=(10, 5))
plt.title('Single Linkage Clustering')
plt.xlabel('Type of Music')
plt.ylabel('Height')
dn = hierarchy.dendrogram(Z, labels=types_list)
```


```



# complete linking 'Type' of music
Z = hierarchy.linkage(standardized_df, 'complete')
plt.figure(figsize=(10, 5))
plt.title('Complete Linkage Clustering')
plt.xlabel('Type of Music')
plt.ylabel('Height')
dn = hierarchy.dendrogram(Z, labels= types_list)


# single linking with Artist labels
Z = hierarchy.linkage(standardized_df, 'single')
plt.figure(figsize=(10, 5))
plt.title('Single Linkage Clustering')
plt.xlabel('Artist')
plt.ylabel('Height')
dn = hierarchy.dendrogram(Z, labels=artists_list)

# complete linking Arist Lables
Z = hierarchy.linkage(standardized_df, 'complete')
plt.figure(figsize=(10, 5))
plt.title('Complete Linkage Clustering')
plt.xlabel('Artist')
plt.ylabel('Height')
dn = hierarchy.dendrogram(Z, labels=artists_list)

```









Which method seems best?

In our opinion, the best method based on the dataset that was provided appears to be the ones that utilize complete linkage. The dendrograms that utilize single linkage appear to have a lot of crowding that makes parsing and interpreting the data more difficult.


Cluster 1




Cluster 2



Cluster 3



Cluster 4



Cluster 5

